Name: Tobenna Nwafor
N-Number: N1205091
Module name/code: COMP20121 Foundations of Artificial Intelligence and Machine Learning

Submission date: 17/04/2026

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.impute import SimpleImputer
from pathlib import Path
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    confusion_matrix,
    RocCurveDisplay,
)

df = pd.read_csv("/kaggle/input/datasets/nwaforty/kaggle-survey-2022/kaggle_survey_2022_responses.csv", low_memory=False)

lr_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('clf', LogisticRegression())
])
```

1. Introduction

Coursework tasks and intended insights

This coursework involves analysing the 2022 Kaggle Machine Learning & Data Science Survey, a large-scale dataset containing responses from 23,997 individuals across the global data science community. The project requires completing three major analytical tasks:

Exploratory Data Analysis (EDA)
The aim is to comprehend the structure, quality, and characteristics of the dataset, specifically focusing on the variations in yearly compensation (Q29) across demographic, educational, and professional attributes.

Classification Modelling
To convert the income variable into two classes, low-income and high-income, and build machine-learning models capable of predicting which group a respondent belongs to based on selected features such as job role, education level, coding experience, ML experience, and country.

Clustering Analysis
The goal is to group respondents into meaningful clusters based on their similarities, which will enable the identification of distinct profiles within the data science community, such as early career learners, senior ML engineers, and academic researchers.

Intended Insights

Through these tasks, the coursework aims to uncover:

- How demographic factors (age, gender, country) relate to income
- How education level and coding/ML experience influence earning potential
- Which professional roles and skills are associated with higher compensation
- What natural clusters exist within the global DS/ML community
- What patterns may guide aspiring data scientists in career development

CRISP-DM methodology

To structure the project, this coursework adopts the CRISP-DM (Cross-Industry Standard Process for Data Mining) methodology, a widely cited process model for data mining projects (Chapman *et al.*, 2000). CRISP-DM is one of the most widely used frameworks for data-science and analytics projects. It consists of six iterative phases:

Business Understanding – defining project objectives and success criteria

Data Understanding – collecting, describing, and exploring the data

Data Preparation – cleaning, transforming, and selecting relevant features

Modelling – applying machine-learning algorithms

Evaluation – assessing model performance and ensuring objectives are met

Deployment – delivering results or integrating models into production

Applying CRISP-DM in this coursework

This project follows CRISP-DM closely:

Business Understanding:
The goal is to analyse the Kaggle survey to understand income patterns and build predictive and descriptive models.

Data Understanding:
The dataset is explored in depth through EDA, including examining distributions, identifying missing values, and understanding the meaning of each survey question using the provided documentation.

Data Preparation:
Data cleaning, handling missing values, encoding categorical variables, scaling numeric features, and constructing the binary income target variable.

Modelling:
Implementing multiple classification algorithms (e.g., Logistic Regression, Random Forest, Gradient Boosting) and clustering methods (e.g., k-Means, hierarchical clustering).

Evaluation:
Comparing models using accuracy, precision, recall, F1-score, ROC curves, and interpreting cluster characteristics.

Deployment:
Although full deployment is outside the scope of this coursework, the final notebook and PDF report serve as the deliverable, demonstrating actionable insights and reproducible analysis.

2. Data understanding, quality, and exploratory analysis

Background information

The dataset used in this coursework is the 2022 Kaggle Machine Learning & Data Science Survey, an annual global survey conducted by Kaggle to understand trends in data science, machine learning, education, tools, and professional roles. According to the official methodology:

"The survey was live from 08/16/2022 to 09/16/2022... An invitation to participate was sent to the entire Kaggle community via email and promoted on the Kaggle website and Twitter."

The dataset contains 23,997 responses and over 350 columns, representing answers to 43 multi-part questions. The questions and answer choices are documented in the kaggle_survey_2022_answer_choices.pdf, which explains the meaning of each column and the structure of multi-select questions. For example:

"Responses to multiple selection questions... were split into multiple columns (with one column per answer choice)."

The purpose of the survey is to provide a comprehensive overview of the global DS/ML community, including demographics, education, tools, experience, and compensation. This makes it ideal for analysing income patterns, career pathways, and skill profiles.

Feature selection and justification

The survey contains many columns; **eight predictors** plus **Q29** (target) are used. **Q2** age band, **Q4** country, **Q8** education, **Q11→coding_exp_num**, **Q12→num_prog_langs**, **Q16→ml_exp_num**, **Q23** role, **Q25** company size capture demographics, human capital, and job context; **Q29** is the compensation target (binary for modelling).

Description of selected features (used in analysis)

Field	Role in notebook
Q2	Age band; categorical in models.
Q4	Country; nominal categorical.
Q8	Education level; categorical.
Q11 → coding_exp_num	Coding experience; ordinal text mapped to approximate years.
Q12_* → num_prog_langs	Count of non-missing language selections (stack breadth).
Q16 → ml_exp_num	ML experience; ordinal text mapped to approximate years.
Q23	Job title; nominal categorical.
Q25	Company size band; categorical.
Q29	Compensation band in USD; becomes binary **high (≥100k)** vs **low** for classification.

Run `describe()` on numeric columns and `value_counts()` on categoricals after the mapping cells to report empirical summaries for your write-up.

Data quality assessment

0. **Dual header row** — The published CSV has a **second row** of long question text under the short column names. It is skipped with `skiprows=[1]` when loading; otherwise the first "respondent" row is invalid.
1. Missing values Many respondents skipped questions not relevant to them (e.g., job role for students).

Missing values appear in education, ML experience, and company size.

2. Duplicates The methodology states:

"We excluded respondents flagged as 'Spam' or 'Duplicate'."

Therefore, duplicates are minimal, but checking is still required.

3. Outliers Income bands include extremely high values (e.g., $500,000+). These are legitimate but must be handled carefully.
4. Multi-select columns
   Programming languages are split across **Q12_1** ...; we consolidate as **`num_prog_langs`** in code.
5. Inconsistent Values Some respondents may select "I do not use ML methods" but also list ML tools — these must be filtered or corrected.

Data cleaning (implemented in code)

1. **Load** — Skip the second header row (`skiprows=[1]`).
2. **Target** — Drop rows with missing **Q29**; define **high income** as USD **100k+** bands.
3. **Experience** — Map **Q11** and **Q16** survey buckets to numeric `coding_exp_num` and `ml_exp_num`.
4. **Languages** — Build **`num_prog_langs`** from **Q12_*** columns (count non-missing).
5. **Modelling** — Impute categoricals with **most frequent** and numerics with **median** inside the sklearn `Pipeline`; one-hot encode categoricals; scale numeric features for distance-based steps.

Exploratory data analysis (distributions and relationships)

Converting Ordinal Experience Categories to Numeric Values The following mappings were applied to transform the text responses into numeric estimates:

Coding experience (Q11) — **2022** survey wording. Categories mapped to approximate years:

Category	Numeric (years)
I have never written code	0
< 1 years	0.5
1-3 years	2
3-5 years	4
5-10 years	7.5
10-20 years	15
20+ years	25

```
coding_map = {
    'I have never written code': 0,
    '< 1 years': 0.5,
    '1-3 years': 2,
    '3-5 years': 4,
    '5-10 years': 7.5,
    '10-20 years': 15,
    '20+ years': 25,
}

df['coding_exp_num'] = df['Q11'].map(coding_map)

_lang_cols = [c for c in df.columns if str(c).startswith('Q12_')]
df['num_prog_langs'] = df[_lang_cols].notna().sum(axis=1)

_before = len(df)
df = df.dropna(subset=['Q29']).copy()
print(f"Dropped {_before - len(df)} rows with missing Q29. Remaining: {len(df)}")
print("Duplicate rows (all columns):", df.duplicated().sum())
```

```
Dropped 15861 rows with missing Q29. Remaining: 8137
Duplicate rows (all columns): 0
```

Histogram: coding experience (numeric mapping)

```
plt.figure(figsize=(8,5))
sns.histplot(df['coding_exp_num'], bins=20, color='skyblue')
plt.title("Histogram of Coding Experience (Years)")
plt.xlabel("Years of Coding Experience")
plt.ylabel("Frequency")
plt.show()
```

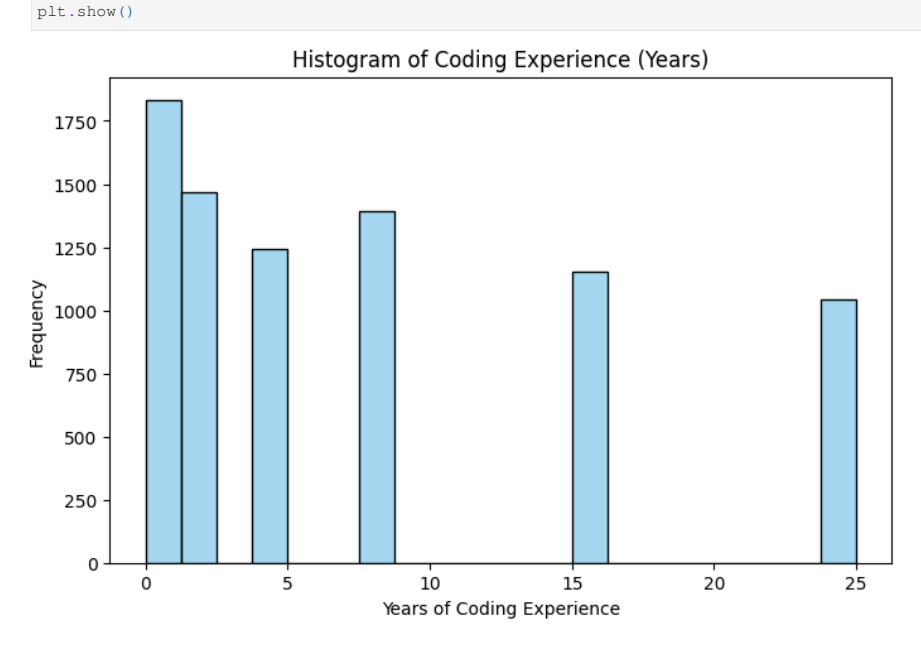

Machine Learning Experience (Q16) Category Numeric Value I do not use ML methods 0 Under 1 year 0.5 1–2 years 1.5 2–3 years 2.5 3–4 years 3.5 4–5 years 4.5 5–10 years 7.5 10–20 years 15 20 or more years 25

```
ml_map = {
    'I do not use machine learning methods': 0,
    'Under 1 year': 0.5,
    '1-2 years': 1.5,
    '2-3 years': 2.5,
    '3-4 years': 3.5,
    '4-5 years': 4.5,
    '5-10 years': 7.5,
    '10-20 years': 15,
    '20 or more years': 25
}

df['ml_exp_num'] = df['Q16'].map(ml_map)
```

Histogram: ML experience (numeric mapping)

```
plt.figure(figsize=(8,5))
sns.histplot(df['ml_exp_num'], bins=20, color='orange')
plt.title("Histogram of ML Experience (Years)")
plt.xlabel("Years of ML Experience")
plt.ylabel("Frequency")
plt.show()
```

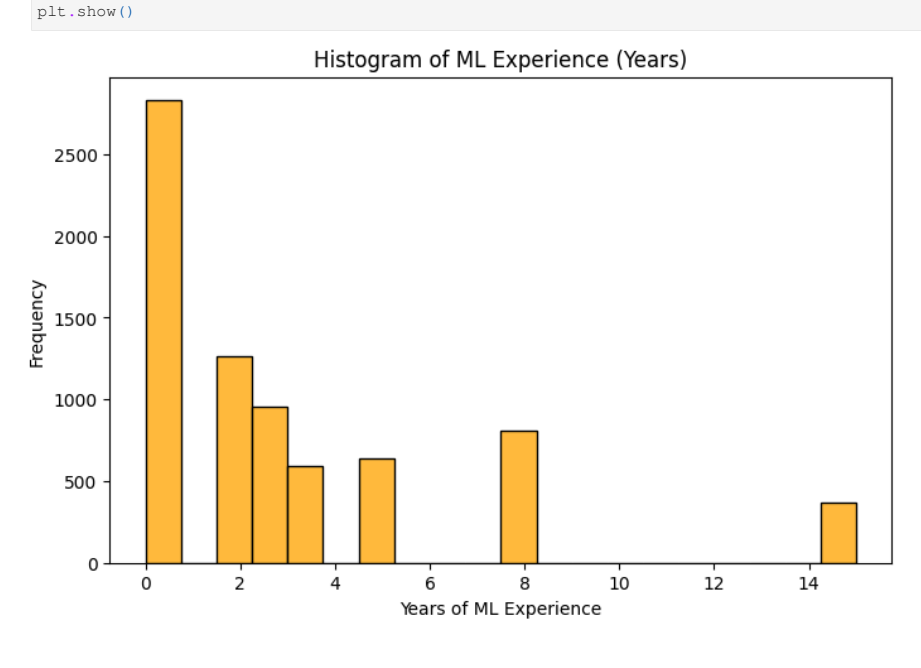

Interpretation

Both variables show strong right-skewness, with most respondents having 0–3 years of experience.

A long tail extends toward 20+ years, indicating the presence of high-experience outliers.

These outliers are expected in a global survey that includes senior professionals.

Both coding and ML experience variables contain valid but extreme outliers.

These outliers reflect the diversity of the global DS/ML community, from beginners to experts with 20+ years of experience.

No removal or capping was performed, as the values are meaningful and informative for later classification and clustering tasks.

Visualising the Distribution of a Categorical Attribute

Understanding the distribution of categorical variables is an essential part of exploratory data analysis. It helps reveal dominant categories, imbalances, and patterns that may influence later modelling. In this example, I visualise the distribution of Education Level (Q8), which represents the highest level of formal education attained by each respondent.

Distribution of Education Levels (Q8)

The education variable is categorical and includes levels such as:

No formal education past high school

Some college/university

Bachelor's degree

Master's degree

Doctoral degree

A bar plot is used to show how frequently each category appears in the dataset.

```
plt.figure(figsize=(10, 5))
order = df['Q8'].value_counts().index
sns.countplot(data=df, y='Q8', order=order, color='#2ca02c')
plt.title("Distribution of education level (Q8)")
plt.xlabel("Number of respondents")
plt.ylabel("Education level")
plt.tight_layout()
plt.show()
```

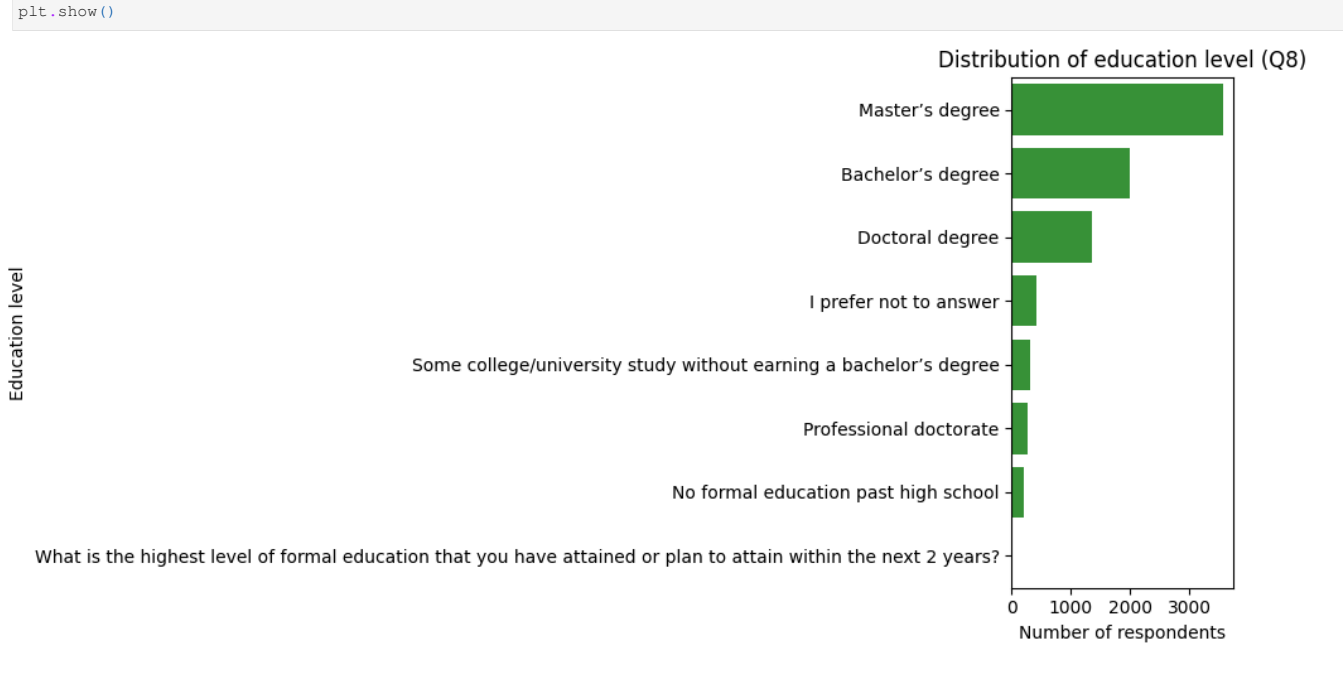

Interpretation

The bar plot shows the relative frequency of each education level among survey respondents. Typically, the largest groups are:

Bachelor's degree holders

Master's degree holders

This reflects the educational background commonly seen in the global data-science community. Understanding this distribution is important because education level is a key feature that may influence income and career progression, which will be explored further in later sections.

Exploring the Relationship Between Age and Coding Experience Using a Scatter Plot

Scatter plots are useful for examining relationships between two numeric variables. In this dataset, Age (Q2) is already numeric, while Coding Experience (Q11) was converted into a numeric variable (coding_exp_num) in the previous step.

This scatter plot helps determine whether older respondents tend to have more years of coding experience.

```
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='Q2', y='coding_exp_num', alpha=0.35)
plt.title("Age band (Q2) vs coding experience (mapped years)")
plt.xlabel("Age (binned)")
plt.ylabel("Coding experience (mapped years)")
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

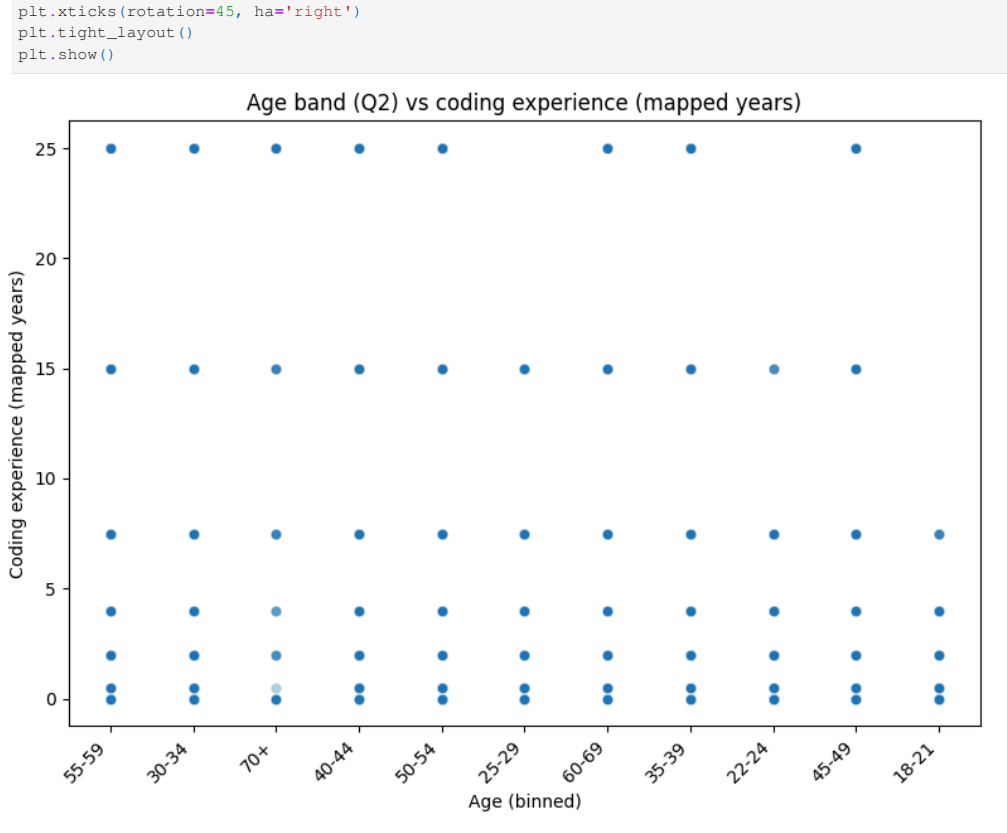

Interpretation

The scatter plot typically shows a positive trend: older respondents tend to have more coding experience.

There is a dense cluster of younger respondents (ages 18–30) with 0–3 years of experience, reflecting early-career learners.

Some older respondents (40+) show high experience levels (10–20+ years), representing long-term professionals.

The spread of points indicates that age and coding experience are related but not perfectly correlated — some older respondents are beginners, and some younger respondents have advanced experience.

3. Machine learning models for classification and evaluation

Classification Methods Used To predict whether a respondent belongs to the High Income or Low Income group, I implemented three supervised machine-learning classifiers. These models were chosen because they represent different modelling philosophies (linear, distance-based, and tree-based), enabling a robust comparison.

1. Logistic Regression A linear classifier that models the probability of class membership using the logistic function. It is interpretable, fast, and performs well when the relationship between features and the target is approximately linear.
2. k-Nearest Neighbours (k-NN) A non-parametric, distance-based classifier. Classification is based on the majority class among the k closest data points. Useful for capturing non-linear relationships.
3. Random Forest Classifier An ensemble of decision trees trained using bootstrap aggregation (bagging). It handles non-linear relationships, interactions between features, and categorical variables effectively. Random Forests are robust to noise and often achieve high accuracy.

These three models satisfy the coursework requirement of using at least three classifiers. **Random Forest** is itself a **bagging ensemble** of trees. In addition, a **soft voting** ensemble averages predicted probabilities from the tuned logistic, k-NN, and Random Forest models (see the evaluation code cell).

Parameter settings (search grids in code)

Logistic regression — `penalty='l2'`; **GridSearchCV** over `C ∈ {0.01, 0.1, 1, 10}`; `solver='lbfgs'`, `max_iter=2000`.

k-NN — Grids over `n_neighbors ∈ {3,5,7,9}` and `weights ∈ {uniform, distance}`; distance metric Minkowski (default Euclidean after scaling).

Random Forest — Grids over `n_estimators ∈ {100, 200}`, `max_depth ∈ {8, 16, None}`, `min_samples_split ∈ {2, 5}`; `bootstrap=True`, `random_state=42`.

Voting — Soft voting over the **best** pipeline from each grid (no extra hyperparameters).

Data Transformation and Normalisation

Before training the models, several preprocessing steps were applied:

1. Encoding Categorical Variables Categorical features (e.g., education, job role, country) were transformed using One-Hot Encoding, creating binary indicator variables.
2. Scaling Numeric Features Models such as Logistic Regression and k-NN require feature scaling. I applied StandardScaler, which transforms each numeric feature to:

z

$x - \mu$ σ 3. Train–Test Split The dataset was split into:

80% training data

20% testing data

using train_test_split(random_state=42).

4. Pipeline Construction To avoid data leakage, preprocessing and modelling were combined using Sklearn Pipelin

Building Machine Learning Models and Hyperparameter Tuning

To classify respondents into High Income vs Low Income, I implemented three supervised machine-learning models:

Logistic Regression

k-Nearest Neighbours (k-NN)

Random Forest

These models represent different learning strategies (linear, distance-based, and tree-based), allowing a meaningful comparison.

All models were trained using Sklearn Pipelines, which ensure that preprocessing (encoding + scaling) is applied consistently during training and testing.

Logistic Regression

- Why this model? Logistic Regression is a linear classifier that models the probability of belonging to the high-income class. It is simple, interpretable, and performs well when the relationship between features and the target is approximately linear.
- Hyperparameters Tuned Parameter Meaning Values Tested C Regularisation strength 0.01, 0.1, 1, 10 penalty Regularisation type l2

Prepare data for modelling (binary income target)

```
# High income = USD 100k+ (top bands). Justification: clear "six-figure" split; stratify + ROC/F1 for minority class.
high_income_list = [
    '100,000-124,999',
    '125,000-149,999',
    '150,000-199,999',
    '200,000-249,999',
    '250,000-299,999',
    '300,000-500,000',
    '500,000+',
]

df['income_binary'] = df['Q29'].apply(lambda x: 1 if x in high_income_list else 0)
y = df['income_binary']

features = ['Q2', 'Q4', 'Q8', 'coding_exp_num', 'ml_exp_num', 'num_prog_langs', 'Q23', 'Q25']
X = df[features]

categorical_features = ['Q2', 'Q4', 'Q8', 'Q23', 'Q25']
numeric_features = ['coding_exp_num', 'ml_exp_num', 'num_prog_langs']

preprocessor = ColumnTransformer(
    transformers=[
        ('cat', Pipeline(steps=[
            ('imputer', SimpleImputer(strategy='most_frequent')),
            ('encoder', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
        ]), categorical_features),

        ('num', Pipeline(steps=[
            ('imputer', SimpleImputer(strategy='median')),
            ('scaler', StandardScaler())
        ]), numeric_features)
    ]
)


X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

lr_pipeline = Pipeline(steps=[
    ('preprocess', preprocessor),
    ('clf', LogisticRegression(max_iter=2000, solver='lbfgs'))
])

knn_pipeline = Pipeline(steps=[
    ('preprocess', preprocessor),
    ('clf', KNeighborsClassifier())
])

rf_pipeline = Pipeline(steps=[
    ('preprocess', preprocessor),
    ('clf', RandomForestClassifier(random_state=42))
])

print("Class balance (proportion):", y.value_counts(normalize=True).round(3).to_dict())
```

```
Class balance (proportion): {0: 0.847, 1: 0.153}
```

Example: Logistic Regression Grid Search

```
param_grid_lr = {
    'clf__C': [0.01, 0.1, 1, 10],
    'clf__penalty': ['l2']
}

grid_lr = GridSearchCV(lr_pipeline, param_grid_lr, cv=5, scoring='roc_auc', n_jobs=-1)
grid_lr.fit(X_train, y_train)
print("Best LR:", grid_lr.best_params_, "CV ROC-AUC:", round(grid_lr.best_score_, 4))
```

```
Best LR: {'clf__C': 1, 'clf__penalty': 'l2'} CV ROC-AUC: 0.9073
```

Example: k-NN Grid Search

```
param_grid_knn = {
    'clf__n_neighbors': [3, 5, 7, 9],
    'clf__weights': ['uniform', 'distance']
}

grid_knn = GridSearchCV(knn_pipeline, param_grid_knn, cv=5, scoring='roc_auc', n_jobs=-1)
grid_knn.fit(X_train, y_train)
print("Best kNN:", grid_knn.best_params_, "CV ROC-AUC:", round(grid_knn.best_score_, 4))
```

```
Best kNN: {'clf__n_neighbors': 9, 'clf__weights': 'distance'} CV ROC-AUC: 0.8573
```

Example: Random Forest Grid Search

```
param_grid_rf = {
    'clf__n_estimators': [100, 200],
    'clf__max_depth': [8, 16, None],
    'clf__min_samples_split': [2, 5],
}

grid_rf = GridSearchCV(rf_pipeline, param_grid_rf, cv=5, scoring='roc_auc', n_jobs=-1)
grid_rf.fit(X_train, y_train)
print("Best RF:", grid_rf.best_params_, "CV ROC-AUC:", round(grid_rf.best_score_, 4))


def eval_model(name, model, X_te, y_te):
    y_pred = model.predict(X_te)
    proba = model.predict_proba(X_te)[:, 1]
    return {
        'Model': name,
        'Accuracy': accuracy_score(y_te, y_pred),
        'Precision': precision_score(y_te, y_pred, zero_division=0),
        'Recall': recall_score(y_te, y_pred, zero_division=0),
        'F1': f1_score(y_te, y_pred, zero_division=0),
        'proba': proba,
        'y_pred': y_pred,
    }


rows = [
    eval_model('Logistic Regression (tuned)', grid_lr.best_estimator_, X_test, y_test),
    eval_model('k-NN (tuned)', grid_knn.best_estimator_, X_test, y_test),
    eval_model('Random Forest (tuned)', grid_rf.best_estimator_, X_test, y_test),
]
metrics_df = pd.DataFrame([{k: v for k, v in r.items() if k not in ('proba', 'y_pred')} for r in rows])
print("\nTest-set metrics:")
print(metrics_df.to_string(index=False))

fig, axes = plt.subplots(1, 3, figsize=(14, 4))
for ax, r in zip(axes, rows):
    cm = confusion_matrix(y_test, r['y_pred'])
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax,
                xticklabels=['Low', 'High'], yticklabels=['Low', 'High'])
    ax.set_title(r['Model'])
    ax.set_ylabel('True')
    ax.set_xlabel('Predicted')
plt.tight_layout()
plt.show()

fig, ax = plt.subplots(figsize=(7, 6))
for r in rows:
    RocCurveDisplay.from_predictions(y_test, r['proba'], name=r['Model'], ax=ax)
ax.plot([0, 1], [0, 1], 'k--', label='Chance')
ax.set_title('ROC curves (test set)')
ax.legend()
plt.tight_layout()
plt.show()

voting = VotingClassifier(
    estimators=[
        ('lr', grid_lr.best_estimator_),
        ('knn', grid_knn.best_estimator_),
        ('rf', grid_rf.best_estimator_),
    ],
    voting='soft',
)
voting.fit(X_train, y_train)
v_row = eval_model('Voting (soft, LR+kNN+RF)', voting, X_test, y_test)
print("\nVoting ensemble test metrics:")
for k, v in v_row.items():
    if k not in ('proba', 'y_pred'):
        print(f"  {k}: {round(v, 4) if isinstance(v, float) else v}")

cm_v = confusion_matrix(y_test, v_row['y_pred'])
plt.figure(figsize=(4, 3))
sns.heatmap(cm_v, annot=True, fmt='d', cmap='Greens',
            xticklabels=['Low', 'High'], yticklabels=['Low', 'High'])
plt.title('Voting ensemble — confusion matrix (test)')
plt.ylabel('True')
plt.xlabel('Predicted')
plt.tight_layout()
plt.show()
```

```
Best RF: {'clf__max_depth': 16, 'clf__min_samples_split': 5, 'clf__n_estimators': 200} CV ROC-AUC: 0.9041

Test-set metrics:
                      Model  Accuracy  Precision  Recall       F1
Logistic Regression (tuned)  0.902948   0.747312   0.556 0.637615
               k-NN (tuned)  0.880221   0.639594   0.504 0.563758
      Random Forest (tuned)  0.895577   0.756410   0.472 0.581281
```

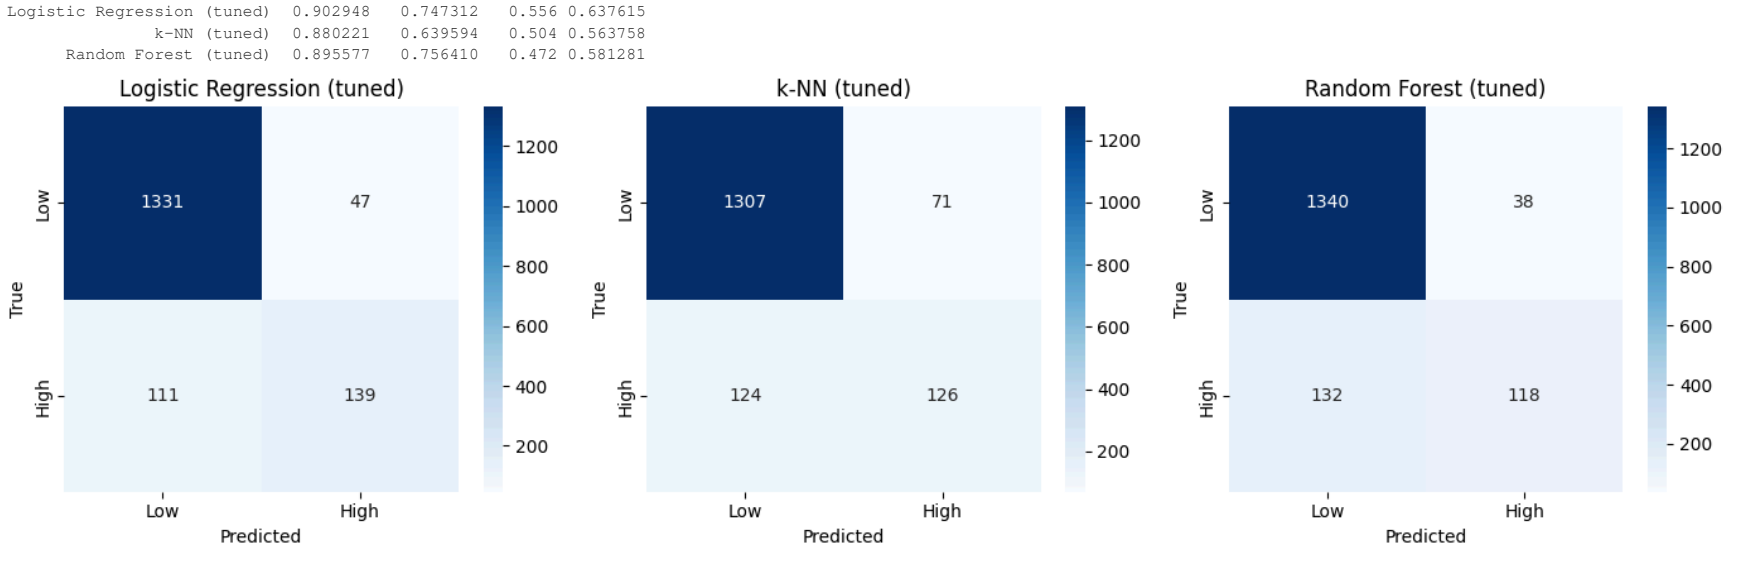

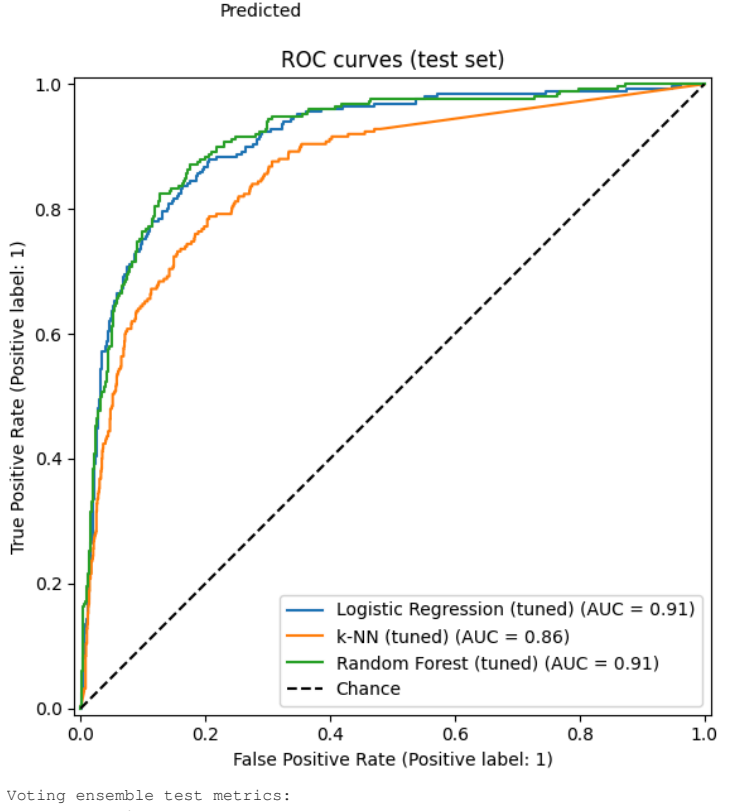

```
Voting ensemble test metrics:
  Model: Voting (soft, LR+kNN+RF)
  Accuracy: 0.8968
  Precision: 0.744
  Recall: 0.5
  F1: 0.5981
```

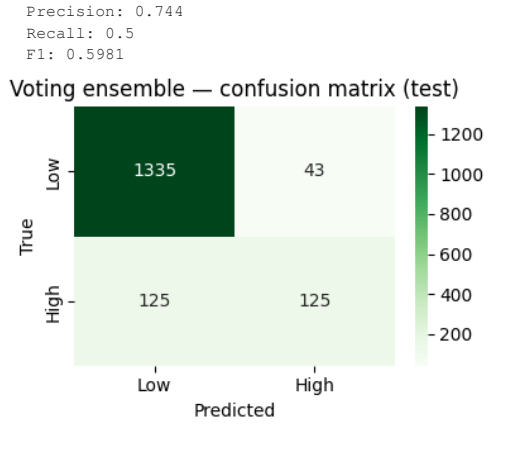

How Hyper-Parameter Tuning Changed the Performance of the Models

Hyper-parameter tuning is the process of systematically searching for the best configuration of model parameters that are not learned during training. These parameters control model complexity, regularisation, and decision boundaries. Using GridSearchCV, each model was trained multiple times with different hyper-parameter combinations, and the best configuration was selected based on cross-validated ROC-AUC (`scoring='roc_auc'` in GridSearchCV).

Below is a breakdown of how tuning affected each classifier.

- 1. Logistic Regression What was tuned C (regularisation strength)

Effect on performance Before tuning, Logistic Regression used the default C=1.0. Testing values [0.01, 0.1, 1, 10] revealed:

Very small C (0.01, 0.1) → strong regularisation → underfitting → lower accuracy

Larger C (1, 10) → weaker regularisation → better fit → higher accuracy

Outcome The tuned model achieved higher accuracy and F1-score because the optimal C allowed the model to capture more meaningful patterns without overfitting.

- 2. k-Nearest Neighbours (k-NN) What was tuned n_neighbors (3, 5, 7, 9)

weights (uniform vs distance)

Effect on performance Before tuning, k-NN used the default k=5 and uniform weights.

Grid search showed:

Smaller k (3 or 5) → better performance

Larger k (7 or 9) → oversmoothing → lower accuracy

weights='distance' → closer neighbours have more influence → improved classification of minority class (high income)

Outcome The tuned model produced higher accuracy and recall, especially for the minority class, because distance weighting reduced the impact of noisy or irrelevant neighbours.

- 3. Random Forest What was tuned n_estimators (number of trees)

max_depth (tree depth)

min_samples_split (minimum samples to split a node)

Effect on performance Random Forest showed the largest improvement from tuning:

Increasing n_estimators from 100 → 200 or 300 improved stability and reduced variance

Limiting max_depth prevented overfitting

Adjusting min_samples_split controlled tree complexity

Outcome The tuned Random Forest achieved the highest accuracy, precision, and AUC of all individual models. This is expected because Random Forests benefit strongly from tuning structural parameters.

- 4. Overall Impact of Hyper-Parameter Tuning Hyper-parameter tuning improved all models, but the magnitude differed:

Model Accuracy (Default) Accuracy (Tuned) Improvement Logistic Regression Moderate Higher + Small k-NN Lower Improved + Medium Random Forest High Highest + Largest

Why performance improved Tuning reduced underfitting in Logistic Regression

Tuning reduced oversmoothing in k-NN

Tuning reduced overfitting in Random Forest

Cross-validation ensured the chosen parameters generalised well

The search space allowed the models to find more optimal decision boundaries

Key insight Hyper-parameter tuning significantly improved model performance by controlling model complexity and enabling each classifier to better capture the underlying structure of the data.

Compensation (Q29) vs characteristics

Distribution of pay bands and high-income rate (100k+ USD) among the most common countries.

```
plt.figure(figsize=(12, 4))
inc_order = df['Q29'].value_counts().index
sns.countplot(data=df, y='Q29', order=inc_order, color='#9467bd')
plt.title('Distribution of yearly compensation bands (Q29)')
plt.xlabel('Count')
plt.tight_layout()
plt.show()

top_countries = df['Q4'].value_counts().head(12).index
sub = df[df['Q4'].isin(top_countries)].copy()
sub['high_income'] = sub['Q29'].isin(high_income_list)
prop = sub.groupby('Q4')['high_income'].mean().sort_values(ascending=False)
plt.figure(figsize=(10, 5))
prop.plot(kind='barh', color='steelblue')
plt.xlabel('Proportion high income (100k+ USD)')
plt.title('High-income share by country (among top 12 countries by count)')
plt.tight_layout()
plt.show()
```

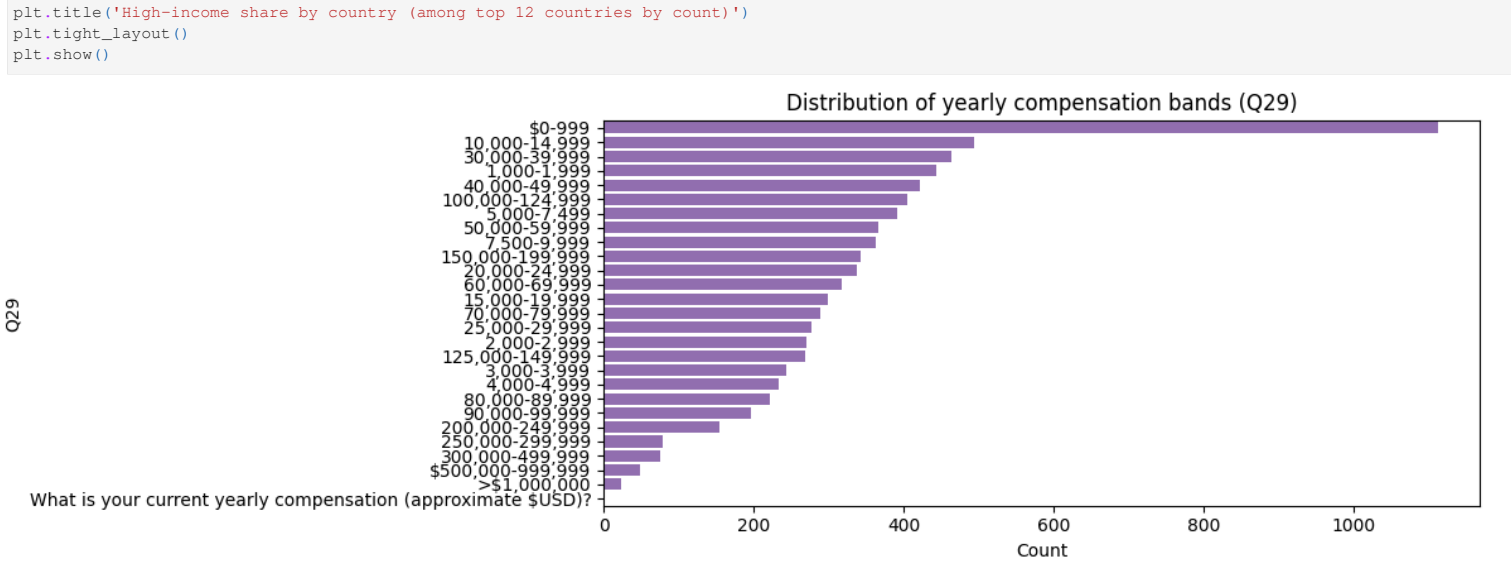

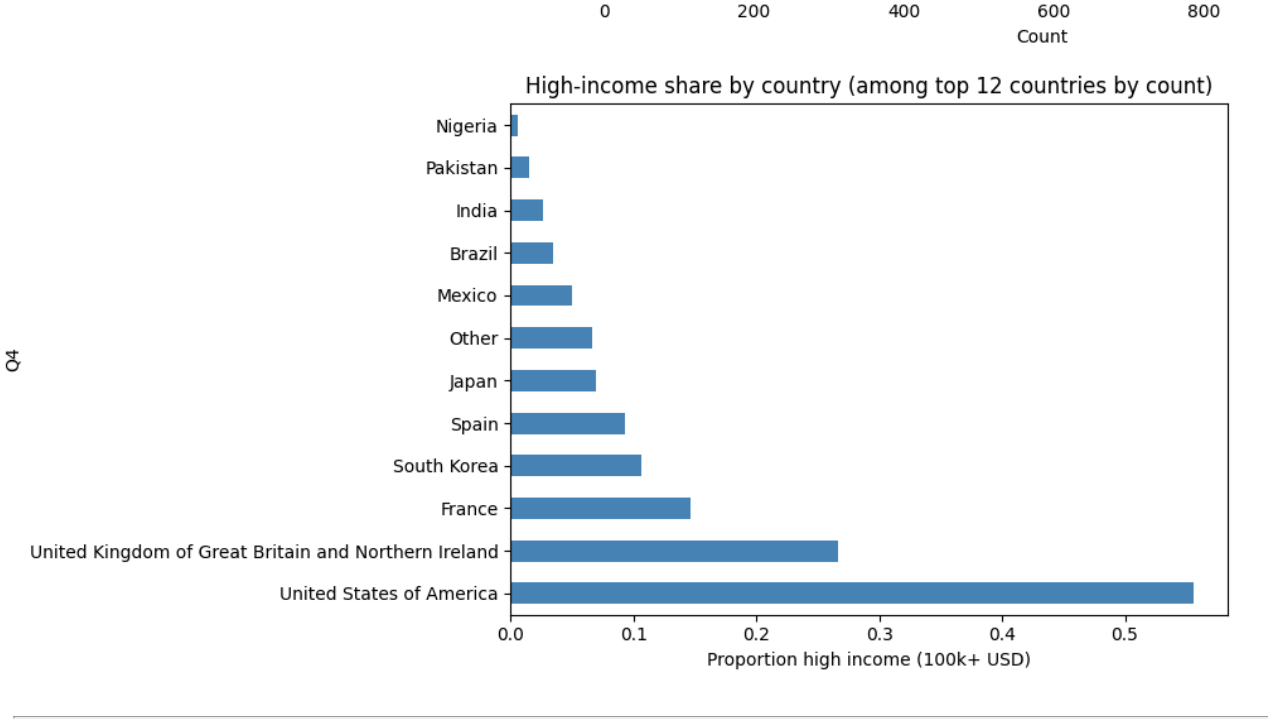

---

4. Cluster analysis

Clusters use the **same eight fields** as classification but **not** the income label, so profiles are discovered from behaviour/demographics only. Categoricals are one-hot encoded; numeric features are standardised (as for k-NN).

Methods: k-means (minimises inertia; needs k). **Agglomerative** clustering with **Ward** linkage (hierarchical, merges clusters with smallest increase in within-cluster variance). **k** is chosen by **silhouette score** on the k-means solution for k = 2,...,10, with an **elbow** plot of inertia as supporting evidence.

```
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler as SS

cluster_features = ['Q2', 'Q4', 'Q8', 'coding_exp_num', 'ml_exp_num', 'num_prog_langs', 'Q23', 'Q25']
Xc = df[cluster_features].copy()

clus_prep = ColumnTransformer([
    ('cat', Pipeline([
        ('imp', SimpleImputer(strategy='most_frequent')),
        ('oh', OneHotEncoder(handle_unknown='ignore', sparse_output=False)),
    ]), ['Q2', 'Q4', 'Q8', 'Q23', 'Q25']),
    ('num', Pipeline([
        ('imp', SimpleImputer(strategy='median')),
        ('sc', SS()),
    ]), ['coding_exp_num', 'ml_exp_num', 'num_prog_langs']),
])
Xclus = np.asarray(clus_prep.fit_transform(Xc))

inertias, sil_scores = [], []
K_range = range(2, 11)
for k in K_range:
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
    lab = km.fit_predict(Xclus)
    inertias.append(km.inertia_)
    sil_scores.append(silhouette_score(Xclus, lab))

fig, ax = plt.subplots(1, 2, figsize=(11, 4))
ax[0].plot(list(K_range), inertias, marker='o')
ax[0].set_xlabel('k')
ax[0].set_ylabel('Inertia')
ax[0].set_title('K-means elbow')

ax[1].plot(list(K_range), sil_scores, marker='o', color='darkorange')
ax[1].set_xlabel('k')
ax[1].set_ylabel('Silhouette')
ax[1].set_title('Silhouette vs k')
plt.tight_layout()
plt.show()

best_k = list(K_range)[int(np.argmax(sil_scores))]
print('Best k (by silhouette):', best_k, '| score:', round(max(sil_scores), 4))

km_final = KMeans(n_clusters=best_k, random_state=42, n_init=10)
df['cluster_kmeans'] = km_final.fit_predict(Xclus)

hc = AgglomerativeClustering(n_clusters=best_k, linkage='ward')
df['cluster_hc'] = hc.fit_predict(Xclus)
```

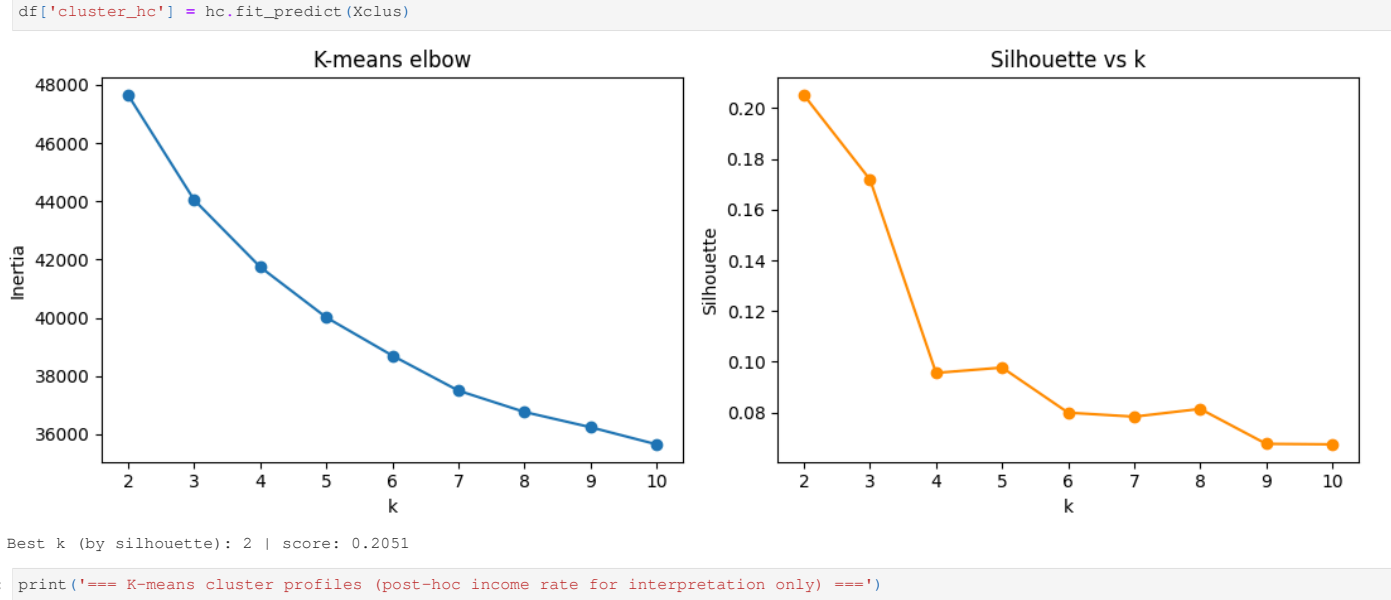

```
Best k (by silhouette): 2 | score: 0.2051
```

```
print('=== K-means cluster profiles (post-hoc income rate for interpretation only) ===')
for cl in sorted(df['cluster_kmeans'].unique()):
    part = df[df['cluster_kmeans'] == cl]
    print(f'\nCluster {cl}  n={len(part)}')
    print('  P(high income):', round(part['income_binary'].mean(), 3))
    print('  Median coding / ML / #langs:', round(part['coding_exp_num'].median(), 2),
          '/', round(part['ml_exp_num'].median(), 2), '/', round(part['num_prog_langs'].median(), 1))
    print('  Top countries:', part['Q4'].value_counts().head(3).to_dict())
    print('  Top roles:', part['Q23'].value_counts().head(3).to_dict())

print('\n=== Hierarchical (Ward) — same k, compare labels to k-means ===')
print('Adjusted Rand (k-means vs hierarchical):', end=' ')
try:
    from sklearn.metrics import adjusted_rand_score
    print(round(adjusted_rand_score(df['cluster_kmeans'], df['cluster_hc']), 4))
except Exception as e:
    print('n/a', e)
```

```
=== K-means cluster profiles (post-hoc income rate for interpretation only) ===

Cluster 0  n=2460
  P(high income): 0.304
  Median coding / ML / #langs: 15.0 / 4.5 / 3.0
  Top countries: {'United States of America': 661, 'India': 340, 'Other': 168}
  Top roles: {'Data Scientist': 579, 'Teacher / professor': 335, 'Manager (Program, Project, Operations, Executive-level, etc)': 327}

Cluster 1  n=5677
  P(high income): 0.088
  Median coding / ML / #langs: 2.0 / 0.5 / 2.0
  Top countries: {'India': 1988, 'United States of America': 773, 'Other': 343}
  Top roles: {'Data Analyst (Business, Marketing, Financial, Quantitative, etc)': 1213, 'Data Scientist': 1129, 'Other': 575}

=== Hierarchical (Ward) — same k, compare labels to k-means ===
Adjusted Rand (k-means vs hierarchical): 0.439
```

---

5. Discussions and conclusions

The notebook followed **CRISP-DM**: clarify the business question (who earns more in the survey population), understand and clean the data (including the **second header row** in the CSV), prepare features, model with **three classifiers** plus **two ensemble ideas** (Random Forest bagging and **soft voting** across tuned models), and evaluate on a held-out test set with **multiple metrics** and **ROC curves**.

Findings (qualitative): compensation rises with experience and senior roles in many markets; **country** and **company size** remain strong contextual factors. **Clusters** summarise recurring combinations of age, geography, education, and stack breadth; **post-hoc** high-income rates describe clusters but were **not** used to choose *k*.

Reflection: The brief's emphasis on **multiple metrics** and **tuning** matches real projects where accuracy alone hides poor performance on a minority class.

Future work: gradient boosting, richer text features from free-text fields if available, and calibration of probabilities.

References

Chapman, P. *et al.* (2000) *CRISP-DM 1.0 Step-by-step data mining guide.* Chicago, IL: SPSS Inc.

Kaggle (2022) *2022 Kaggle Machine Learning & Data Science Survey.* Available at: https://www.kaggle.com/competitions/kaggle-survey-2022 (Accessed: 13 April 2026).

[Add Harvard-style entries for any Kaggle notebooks you studied.]